



Face Recognition: Tech. that knows you!

“All, all are gone, the old familiar faces.”

— Charles Lamb



Face Recognition

The ability to recognize a person's face in a digital image.

Early Days

- First real system was created in 1964 by Woody Bledsoe, Helen Chan Wolf and Charles Bisson.
- Operator had to manually measure and locate key points on the face.
- System could only filter down a list pictures to a small set of possible matches.



GPUs and NNs revolt

- Systems using increasingly complex statistical models were built from the 1970s through the 2000s.
- The availability of GPUs and the advent of deep learning through the mid 2010s.




Wild Apps

- Key research papers from Google and Facebook in 2014 and 2015 demonstrated how to reach near human-level performance on face recognition.
- Available to almost anyone for all kinds of products.




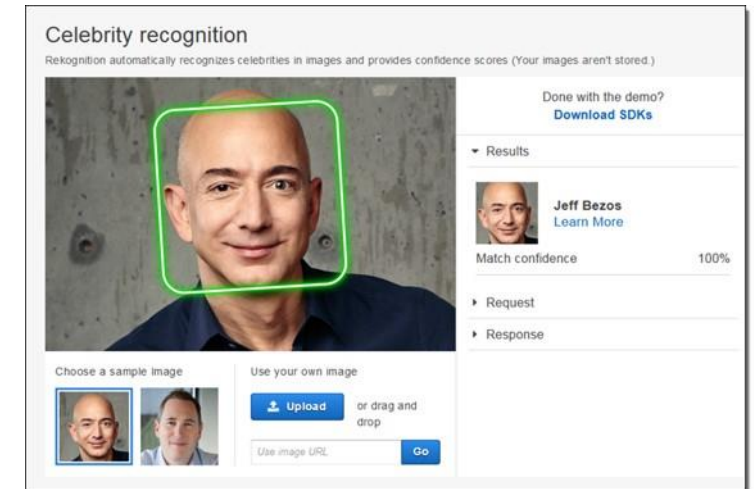
Uses

- Identity Verification
- Tagging Photos
- Surveillance
- Advertisement




Tools

- Commercial Services
 - Amazon Rekognition
 - Microsoft Azure Face
 - Google Cloud Vision
 - IBM Watson Visual Recognition
- Open-Source
 - OpenFace by Brandon Amos at Carnegie Mellon.
 - Dlib by Davis King.



Pareidolia



Pareidolia



Multi-Step Pipeline


- Locate and extract faces from the image.
- Identify face landmarks.
- Align faces to match a pose template.
- Encode faces.
- Check the Euclidean distance between face encodings.



Face Detection

The ability to detect and locate human faces in an image.


Step 1: Build a face detector.



Step 2: Slide a window.



Step 3: Return location.




Face Classification

- Viola-Jones
- Histogram of Gradients or HOG
- Convolutional Neural Networks or CNNs


Viola-Jones

- Invented by Paul Viola and Michael Jones in the early 2000s.
- Very fast and great for low-powered devices.




HOG: Conversion


Input Image



HOG Representation




Step 1: Build a HOG face detector.



**HOG Face
Detector**



True



**HOG Face
Detector**




False

Step 2: Slide a window.




Step 3: Return location.



HOG: Advantages

- Compactness with enough detail to detect faces.
- Resistance to small changes in lighting and object's shape.




CNNs

- Most accurate solution.
- Lots of training data.
- Slow unless backed by a GPU to offload processing.


Face Landmark Estimation

Locating key points on a face, like the center of the eyes and mouth.



Face Landmark Estimation





Face Landmark Estimation







Trick 1: Assume all faces are similar.



Trick 2: Limit movement of each point.



Trick 3: Fine-tune with multiple models.





Face Alignment

Adjusting a raw face image so that key facial features like the eyes, nose and mouth line up with a predefined template.

Step 1: Detect face landmarks.



Step 2: Calculate the Affine Transform.



Affine Transformation

A linear mapping between sets of points where parallel lines remain parallel.

Step 2: Calculate the Affine Transform.



Identifying an Unknown Face



Unknown Face



Known Faces

Comparing Faces





Face Encoding

The process of taking a face image and turning it into a set of measurements.

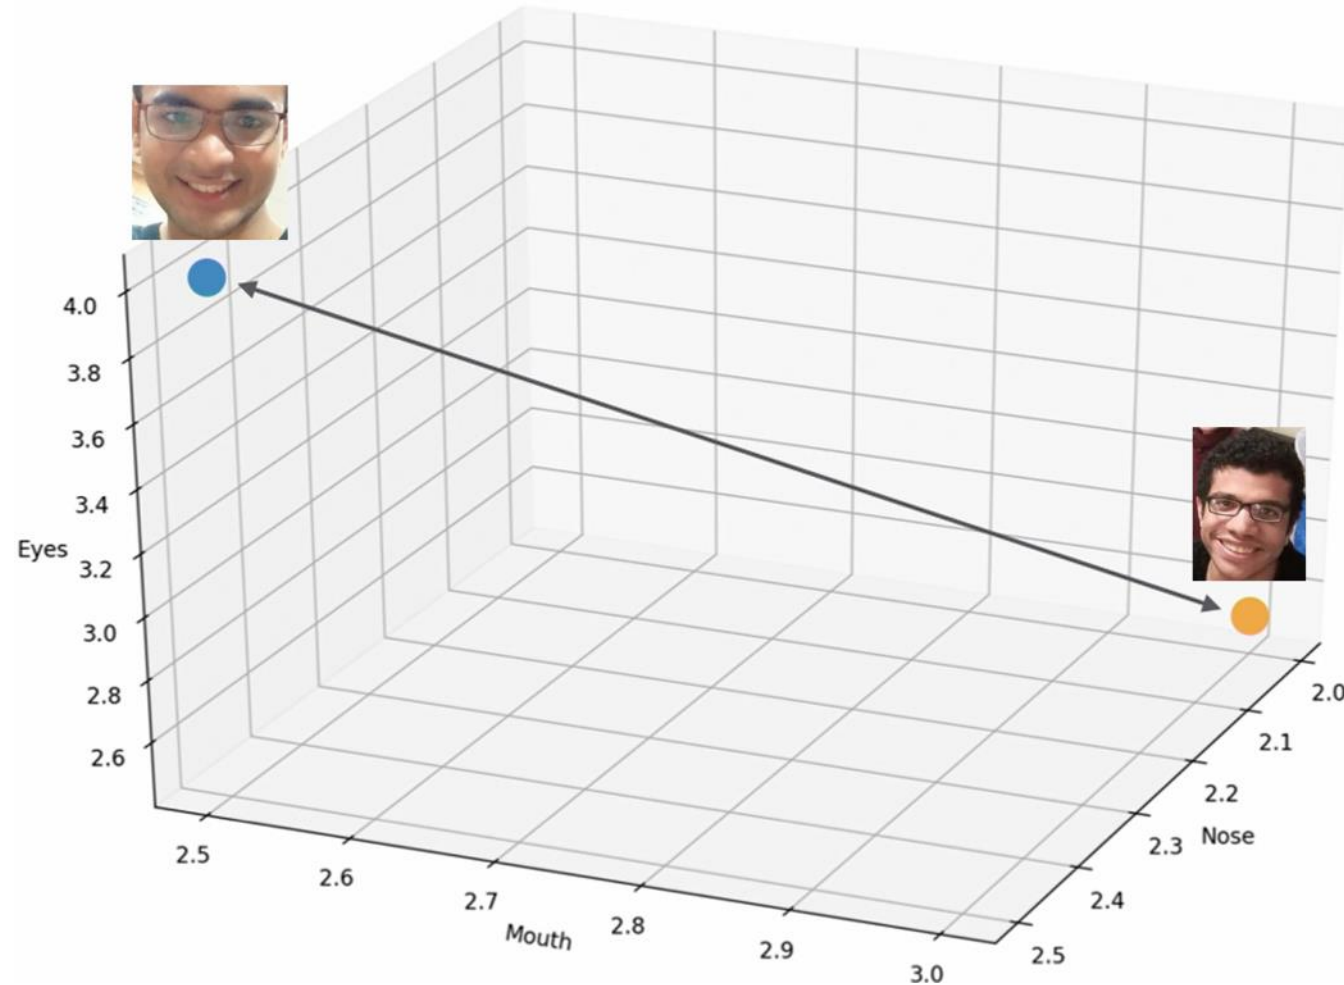
Measuring Faces



	Left	Right
Nose	2.5	2
Mouth	2.5	3
Eyes	4	2.5



Reduced Faces to Points in 3D Space





Euclidean Distance

The distance between two points in space along a straight line.

Face Distance Threshold

- Maximum face distance at which the face is considered the same.
- If $distance(Left, Right) > 0.6$ then they're not a match.
- If $distance(Left, Right) \leq 0.6$ then they're a match.
- The lower the distance, the better the match.

Calculating Euclidean Distance

$$\text{distance}(a, b) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2 + (a_3 - b_3)^2}$$

$$\text{distance}(\text{Left}, \text{Right}) = \sqrt{(2.5 - 2)^2 + (2.5 - 3)^2 + (4 - 2.5)^2}$$

$$\text{distance}(\text{Left}, \text{Right}) = \sqrt{2.75}$$

$$\text{distance}(\text{Left}, \text{Right}) = 1.66 \quad \text{Not a match!}$$

Deep Metric Learning

Using deep learning to have a computer come up with a way to measure something.

Training with Triplets

A



B



Face distance (A, B) should be small.

C




Face distance (A, C) should be large.

Training with Triplets

A



B



Face distance (A, B) should be small.

C



Face distance (A, C) should be large.

Encoding Faces with 128 Points



**Face
Encoder**



-0.6948713
0.11813076
0.22286303
0.00477135
-0.0521124
-0.02841697
0.00964657
0.03243792
0.03346656
0.00654369
0.24236825
-0.14097737
-0.31652586
-0.09467196
... etc ...

Calculating Euclidean Distance with 128 Points

$$distance(\mathbf{a}, \mathbf{b}) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots + (a_{128} - b_{128})^2}$$

Limitations

- Data and Privacy
- Performance
- Security

Demo 😊